

R Notebook

The following is your first chunk to start with. Remember, you can add chunks using the menu above (Insert -> R) or using the keyboard shortcut Ctrl+Alt+I. A good practice is to use different code chunks to answer different questions. You can delete this comment if you like.

Other useful keyboard shortcuts include Alt- for the assignment operator, and Ctrl+Shift+M for the pipe operator. You can delete these reminders if you don't want them in your report.

1.a.

```
#setwd("C:/") #Don't forget to set your working directory before you start!
```

```
library("tidyverse")
```

```
## -- Attaching packages ----- tidyverse  
1.3.0 --
```

```
## v ggplot2 3.2.1      v purrr  0.3.3  
## v tibble  2.1.3      v dplyr  0.8.3  
## v tidyr   1.0.0      v stringr 1.4.0  
## v readr   1.3.1      v forcats 0.4.0
```

```
## -- Conflicts -----  
tidyverse_conflicts() --  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag()     masks stats::lag()
```

```
library("tidymodels")
```

```
## Registered S3 method overwritten by 'xts':  
##   method      from  
##   as.zoo.xts  zoo
```

```
## -- Attaching packages ----- tidymodels  
0.0.3 --
```

```
## v broom      0.5.3      v recipes  0.1.9  
## v dials      0.0.4      v rsample  0.0.5  
## v infer      0.5.1      v yardstick 0.0.4  
## v parsnip    0.0.5
```

```
## -- Conflicts -----  
tidymodels_conflicts() --  
## x scales::discard() masks purrr::discard()  
## x dplyr::filter()   masks stats::filter()  
## x recipes::fixed()  masks stringr::fixed()
```

```

## x dplyr::lag()          masks stats::lag()
## x dials::margin()      masks ggplot2::margin()
## x yardstick::spec()    masks readr::spec()
## x recipes::step()      masks stats::step()
## x recipes::yj_trans()  masks scales::yj_trans()

library("plotly")

##
## Attaching package: 'plotly'

## The following object is masked from 'package:ggplot2':
##
##   last_plot

## The following object is masked from 'package:stats':
##
##   filter

## The following object is masked from 'package:graphics':
##
##   layout

library("skimr")
library("caret")

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following objects are masked from 'package:yardstick':
##
##   precision, recall

## The following object is masked from 'package:purrr':
##
##   lift

dfc<- read_csv("assignment3Carvana.csv")

## Parsed with column specification:
## cols(
##   Auction = col_character(),
##   Age = col_double(),
##   Make = col_character(),
##   Color = col_character(),
##   WheelType = col_character(),
##   Odo = col_double(),
##   Size = col_character(),
##   MMRAuction = col_double(),
##   MMRAretail = col_double(),

```

```
## BadBuy = col_double()
## )
skim(dfc)
```

Data summary

```
Name          dfc
Number of rows 10061
Number of columns 10
```

Column type frequency:






```
character      5
numeric        5
```

```
Group variables      None
```

Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
Auction	0	1	5	7	0	3	0
Make	0	1	3	10	0	30	0
Color	0	1	3	8	0	17	0
WheelType	0	1	4	7	0	4	0
Size	0	1	3	10	0	12	0


Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
Age	0	1	4.50	1.77	1	3	4	6	9	
Odo	0	1	72903.87	14498.87	94	634	749	836	1157	
MMRAuction	0	1	5812.38	2578.85	0	387	558	745	3572	
MMRAetail	0	1	8171.51	3257.19	0	587	805	103	3908	
BadBuy	0	1	0.50	0.50	0	0	0	1	1	

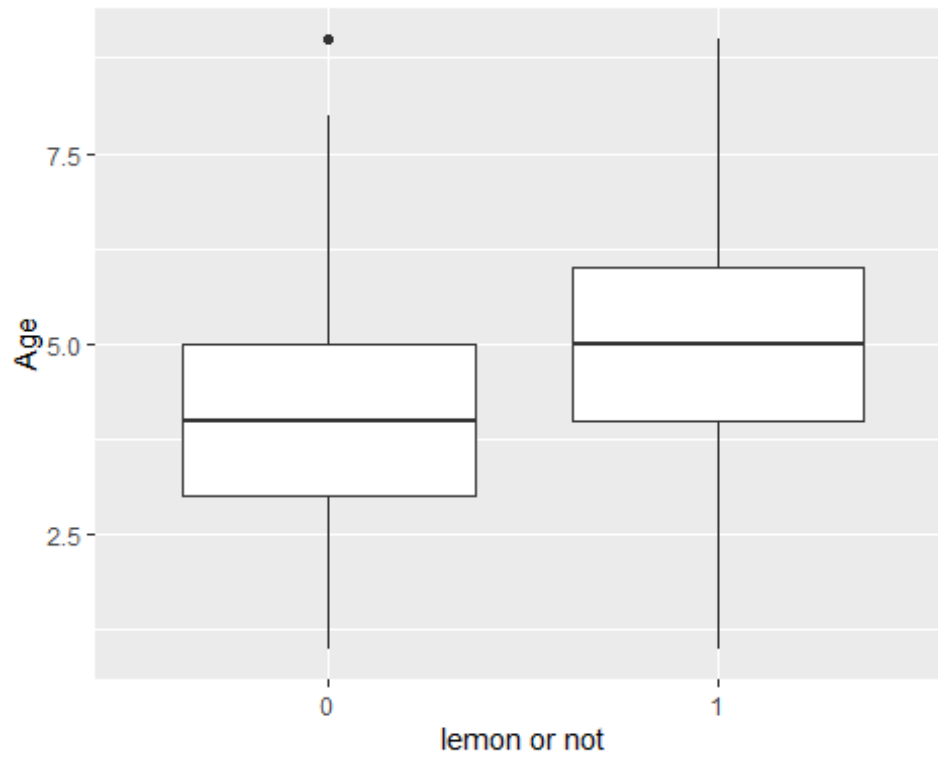
1.b.

```
set.seed(52156)
dfcTrain1<- dfc %>% sample_frac(0.65)
dfcTest1<- dplyr::setdiff(dfc, dfcTrain1)


plot1<- dfcTrain1%>%
  ggplot()+geom_boxplot(aes(x=as.factor(BadBuy), y=MMRAuction))+xlab("lemon
or not")
plot1
```



```
plot2<- dfcTrain1%>%
  ggplot()+geom_boxplot(aes(x=as.factor(BadBuy), y=Age))+xlab("lemon or not")
plot2
```



```
plot3<- dfcTrain1%>%
  ggplot()+geom_boxplot(aes(x=as.factor(BadBuy), y=Odo))+xlab("lemon or not")
plot3
```



2.b.

```
dfcTrain1 %>%
  group_by(BadBuy, Size) %>%
  tally() %>%
  mutate(pct = 100*n/sum(n)) %>%
  arrange(desc(BadBuy), desc(pct))

## # A tibble: 24 x 4
## # Groups:   BadBuy [2]
##   BadBuy Size          n    pct
##   <dbl> <chr>      <int> <dbl>
## 1     1  MEDIUM    1298  39.8
## 2     1  COMPACT     448  13.7
## 3     1  MEDIUMSUV   412  12.6
## 4     1  LARGE       284   8.70
## 5     1  VAN         269   8.24
## 6     1  LARGETRUCK  126   3.86
## 7     1  SMALLSUV    112   3.43
## 8     1  LARGESUV     76   2.33
## 9     1  SPECIALTY    68   2.08
## 10    1  CROSSOVER    66   2.02
## # ... with 14 more rows

fitem<-
  lm(formula = BadBuy~ . , data=dfcTrain1)

summary(fitem)

##
## Call:
## lm(formula = BadBuy ~ ., data = dfcTrain1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.2353 -0.3934 -0.1635  0.4658  0.9587
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -1.996e-01  2.394e-01  -0.834  0.40434
## AuctionMANHEIM  4.065e-02  1.490e-02   2.728  0.00638 **
## AuctionOTHER    2.287e-02  1.706e-02   1.341  0.18008
## Age            5.154e-02  5.619e-03   9.172 < 2e-16 ***
## MakeBUICK      2.392e-01  2.360e-01   1.013  0.31089
## MakeCADILLAC   2.664e-01  5.045e-01   0.528  0.59756
## MakeCHEVROLET  1.861e-01  2.299e-01   0.810  0.41820
## MakeCHRYSLER   2.944e-01  2.297e-01   1.282  0.19993
## MakeDODGE      2.384e-01  2.293e-01   1.040  0.29853
## MakeFORD       2.620e-01  2.298e-01   1.140  0.25427
## MakeGMC        1.398e-01  2.379e-01   0.588  0.55685
## MakeHONDA      1.114e-01  2.374e-01   0.469  0.63904
```

## MakeHYUNDAI	2.099e-01	2.321e-01	0.904	0.36578	
## MakeINFINITI	3.671e-01	3.201e-01	1.147	0.25141	
## MakeISUZU	1.764e-01	2.747e-01	0.642	0.52082	
## MakeJEEP	2.537e-01	2.331e-01	1.089	0.27638	
## MakeKIA	2.190e-01	2.316e-01	0.946	0.34440	
## MakeLEXUS	8.805e-01	3.221e-01	2.733	0.00629	**
## MakeLINCOLN	3.712e-01	2.577e-01	1.440	0.14980	
## MakeMAZDA	2.567e-01	2.329e-01	1.102	0.27036	
## MakeMERCURY	2.980e-01	2.337e-01	1.275	0.20229	
## MakeMINI	3.301e-01	3.082e-01	1.071	0.28422	
## MakeMITSUBISHI	1.179e-01	2.338e-01	0.504	0.61396	
## MakeNISSAN	2.310e-01	2.313e-01	0.999	0.31801	
## MakeOLDSMOBILE	3.261e-01	2.441e-01	1.336	0.18156	
## MakePONTIAC	2.181e-01	2.306e-01	0.946	0.34427	
## MakeSATURN	2.800e-01	2.316e-01	1.209	0.22684	
## MakeSCION	1.091e-01	2.669e-01	0.409	0.68272	
## MakeSUBARU	2.432e-01	3.922e-01	0.620	0.53520	
## MakeSUZUKI	3.696e-01	2.335e-01	1.583	0.11354	
## MakeTOYOTA	1.638e-01	2.341e-01	0.700	0.48414	
## MakeVOLKSWAGEN	2.630e-01	2.613e-01	1.007	0.31409	
## MakeVOLVO	-1.809e-01	3.906e-01	-0.463	0.64322	
## ColorBLACK	2.220e-02	4.160e-02	0.534	0.59365	
## ColorBLUE	1.890e-02	4.055e-02	0.466	0.64111	
## ColorBROWN	1.819e-02	7.917e-02	0.230	0.81826	
## ColorGOLD	5.438e-02	4.271e-02	1.273	0.20298	
## ColorGREEN	2.264e-02	4.620e-02	0.490	0.62408	
## ColorGREY	3.804e-02	4.137e-02	0.919	0.35793	
## ColorMAROON	7.248e-02	5.097e-02	1.422	0.15503	
## ColorNOTAVAIL	-4.753e-02	1.265e-01	-0.376	0.70717	
## ColorNULL	-1.179e-01	4.546e-01	-0.259	0.79543	
## ColorORANGE	4.598e-02	8.977e-02	0.512	0.60852	
## ColorOTHER	-1.388e-01	9.958e-02	-1.394	0.16327	
## ColorPURPLE	1.955e-02	8.259e-02	0.237	0.81289	
## ColorRED	6.169e-02	4.214e-02	1.464	0.14326	
## ColorSILVER	4.814e-02	3.960e-02	1.216	0.22418	
## ColorWHITE	6.047e-02	4.013e-02	1.507	0.13186	
## ColorYELLOW	-6.072e-02	1.016e-01	-0.597	0.55031	
## WheelTypeCovers	-3.534e-02	1.395e-02	-2.533	0.01134	*
## WheelTypeNULL	5.096e-01	1.861e-02	27.379	< 2e-16	***
## WheelTypeSpecial	-8.805e-03	5.743e-02	-0.153	0.87815	
## Odo	2.888e-06	4.327e-07	6.675	2.69e-11	***
## SizeCROSSOVER	-1.783e-01	4.404e-02	-4.048	5.23e-05	***
## SizeLARGE	-1.475e-01	2.616e-02	-5.640	1.77e-08	***
## SizeLARGESUV	-1.379e-01	4.893e-02	-2.817	0.00486	**
## SizeLARGETRUCK	-1.916e-01	3.669e-02	-5.224	1.81e-07	***
## SizeMEDIUM	-9.926e-02	2.020e-02	-4.913	9.18e-07	***
## SizeMEDIUMSUV	-9.874e-02	2.840e-02	-3.477	0.00051	***
## SizeSMALLSUV	-1.333e-01	4.231e-02	-3.149	0.00164	**
## SizeSMALLTRUCK	-1.449e-01	5.170e-02	-2.803	0.00508	**
## SizeSPECIALTY	-7.220e-02	4.718e-02	-1.530	0.12599	

```
## SizeSPORTS      -1.081e-01  5.064e-02  -2.135  0.03277 *
## SizeVAN         -1.136e-01  2.727e-02  -4.164  3.16e-05 ***
## MMRAauction     1.595e-06  7.264e-06   0.220  0.82626
## MMRAretail      -1.126e-06  4.514e-06  -0.249  0.80302
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4502 on 6474 degrees of freedom
## Multiple R-squared:  0.1975, Adjusted R-squared:  0.1894
## F-statistic: 24.51 on 65 and 6474 DF,  p-value: < 2.2e-16
```

3.a

```
testresult <- dfcTest1 %>%
  mutate(predictedtest = predict(fitem, dfcTest1))
testresult

## # A tibble: 3,521 x 11
##   Auction Age Make Color WheelType Odo Size MMRAauction MMRAretail
##   <chr>   <dbl> <chr> <chr> <chr>   <dbl> <chr>   <dbl>   <dbl>
##   <dbl>
## 1 MANHEIM      6 SATU~ WHITE Covers    81116 MEDI~    2667    3380
## 2 OTHER        5 CHEV~ RED Alloy    54718 MEDI~    6921    7975
## 3 OTHER        5 CHEV~ GOLD Covers    89365 VAN    6131    9793
## 4 ADESA        3 CHEV~ WHITE Covers    71794 VAN    6394    7406
## 5 OTHER        3 CHEV~ WHITE NULL    67229 COMP~    5785    9834
## 6 MANHEIM      3 DODGE GOLD Covers    71079 MEDI~    4297    5141
## 7 MANHEIM      6 OLDS~ SILV~ Alloy    71235 MEDI~    3325    4091
## 8 MANHEIM      8 PONT~ SILV~ Alloy    90325 MEDI~    2150    4937
## 9 MANHEIM      6 PONT~ GREEN Alloy    96893 MEDI~    4059    4884
## 10 OTHER       2 DODGE BLUE Covers    45151 MEDI~    7982    9121
## # ... with 3,511 more rows, and 1 more variable: predictedtest <dbl>

performance <-
  metric_set(rmse, mae)
performance(testresult, truth= BadBuy, estimate = predictedtest)

## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
```



```
## 1 rmse      standard      0.453
## 2 mae       standard      0.415

trainresult <- dfcTrain1 %>%
  mutate(predictedtest = predict(fitem, dfcTrain1))
trainresult

## # A tibble: 6,540 x 11
##   Auction Age Make Color WheelType Odo Size MMRAauction MMRAretail
BadBuy
##   <chr>   <dbl> <chr> <chr> <chr>   <dbl> <chr>   <dbl>   <dbl>
<dbl>
## 1 MANHEIM     4 FORD  SILV~ NULL    77591 LARG~    9774    14506
1
## 2 MANHEIM     5 MINI  BLUE  Alloy   80013 COMP~   11040    12423
1
## 3 MANHEIM     2 CHEV~ SILV~ Covers  75493 LARGE    9707    13975
1
## 4 ADESA       4 NISS~ BLUE  NULL    84827 MEDI~    6073     9791
1
## 5 MANHEIM     5 FORD  GREY  Alloy   57388 SPOR~    5574     8984
1
## 6 ADESA       4 SUZU~ BLACK NULL    75822 MEDI~    4033     6979
1
## 7 MANHEIM     2 KIA   BLACK Covers  51059 MEDI~    4839     5726
0
## 8 OTHER       7 FORD  GREY  NULL    74595 LARG~    7649    11059
1
## 9 MANHEIM     6 FORD  BLUE  Alloy   80328 LARG~    6172     7166
0
## 10 MANHEIM    8 PONT~ WHITE Alloy   97173 LARGE    3242     6225
0
## # ... with 6,530 more rows, and 1 more variable: predictedtest <dbl>

performance <-
  metric_set(rmse, mae)
performance(trainresult, truth= BadBuy, estimate = predictedtest)


## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 rmse    standard      0.448
## 2 mae     standard      0.410
```

3.c.

```
lemresults <-
  lm(formula = BadBuy ~ . , data = dfcTrain1) %>%
  predict(dfcTest1, type = 'response') %>%
  bind_cols(dfcTest1, predictedProb=.) %>%
  mutate(predictedclass = as.factor(ifelse(predictedProb>0.5,1,0)))
```

```
lemresults$BadBuy<- as.factor(lemresults$BadBuy)

lemresults %>%
  conf_mat(truth = (BadBuy) ,estimate = predictedclass) %>%
  autoplot(type = 'heatmap')
```



```
result = data.frame(Auction="ADESA", Age=1, Make="HONDA",Color="SILVER",
WheelType="Covers",Odo=10000, Size="LARGE",MMRAuction=8000,
MMRAREtail=10000)
```

```
predict(fitlem, result, type="response")
```

```
##          1
## -0.1410712
```

4.a.

```
dfcTrain1<- dfcTrain1%>%
  mutate(BadBuy=as.factor(BadBuy))
dfcTest1<- dfcTest1%>%
  mutate(BadBuy=as.factor(BadBuy))

fitlemon <-
  train(BadBuy~ ., data= dfcTrain1, family='binomial', method='glm') %>%
  predict(dfcTest1 , type='raw')%>%
  bind_cols(dfcTest1, predictedProb=.)
```

[illegible]

```

== :
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type
== :
## prediction from a rank-deficient fit may be misleading

```

4.a.i

```

dfc <- dfc%>%
  mutate(Color = if_else(Color == "NULL", "NOTAVAIL", Color))

dfc<- dfc%>%
  mutate(Make = if_else(Make %in% c("ACURA", "CADILLAC", "VOLVO",
"SUBARU", "MINI", "LEXUS", NA), "OTHER", Make))

dfc$BadBuy <- as.factor(dfc$BadBuy)

set.seed(52156)
dfcTrain1<- dfc %>% sample_frac(0.65)
dfcTest1<- dplyr::setdiff(dfc, dfcTrain1)

dfcTrain1<- dfcTrain1%>%
  mutate(BadBuy=as.factor(BadBuy))
dfcTest1<- dfcTest1%>%
  mutate(BadBuy=as.factor(BadBuy))

fitlemon <-
  glm(BadBuy~ ., data= dfcTrain1, family='binomial')

summary(fitlemon)

```


```
##
## Call:
## glm(formula = BadBuy ~ ., family = "binomial", data = dfcTrain1)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.0725  -0.9782  -0.4717   1.0946   2.1705
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -2.472e+00  4.513e-01  -5.478 4.30e-08 ***
## AuctionMANHEIM  1.735e-01  7.493e-02   2.316 0.020579 *
## AuctionOTHER    9.519e-02  9.037e-02   1.053 0.292217
## Age            2.785e-01  2.887e-02   9.647 < 2e-16 ***
## MakeCHEVROLET  -2.774e-01  2.895e-01  -0.958 0.337982
## MakeCHRYSLER    2.527e-01  3.011e-01   0.839 0.401419
## MakeDODGE       -2.483e-02  2.966e-01  -0.084 0.933287
## MakeFORD         1.020e-01  2.945e-01   0.346 0.729155
## MakeGMC          -5.054e-01  4.193e-01  -1.205 0.228054
## MakeHONDA        -6.530e-01  4.317e-01  -1.512 0.130433
## MakeHYUNDAI     -1.623e-01  3.381e-01  -0.480 0.631275
## MakeINFINITI     3.727e-01  1.280e+00   0.291 0.771007
## MakeISUZU       -3.227e-01  7.887e-01  -0.409 0.682408
## MakeJEEP         3.121e-02  3.496e-01   0.089 0.928850
## MakeKIA          -9.342e-02  3.281e-01  -0.285 0.775823
## MakeLINCOLN      6.866e-01  7.410e-01   0.927 0.354146
## MakeMAZDA        3.015e-02  3.530e-01   0.085 0.931925
## MakeMERCURY      2.670e-01  3.632e-01   0.735 0.462313
## MakeMITSUBISHI  -6.722e-01  3.692e-01  -1.821 0.068664 .
## MakeNISSAN       -7.824e-02  3.213e-01  -0.243 0.807645
## MakeOLDSMOBILE   4.725e-01  5.397e-01   0.875 0.381344
## MakeOTHER        3.109e-01  6.256e-01   0.497 0.619240
## MakePONTIAC     -1.156e-01  3.039e-01  -0.380 0.703748
## MakeSATURN       2.040e-01  3.293e-01   0.620 0.535513
## MakeSCION        -6.429e-01  7.485e-01  -0.859 0.390426
## MakeSUZUKI       6.756e-01  3.578e-01   1.888 0.058974 .
## MakeTOYOTA      -4.609e-01  3.718e-01  -1.240 0.215081
## MakeVOLKSWAGEN   3.278e-02  6.815e-01   0.048 0.961638
## ColorBLACK       1.502e-01  2.157e-01   0.696 0.486312
## ColorBLUE        1.197e-01  2.103e-01   0.569 0.569124
## ColorBROWN       1.348e-01  3.891e-01   0.346 0.729074
## ColorGOLD         3.066e-01  2.201e-01   1.393 0.163652
## ColorGREEN        1.723e-01  2.369e-01   0.727 0.466976
## ColorGREY         2.307e-01  2.139e-01   1.078 0.280903
## ColorMAROON       4.114e-01  2.596e-01   1.585 0.112963
## ColorNOTAVAIL    -2.898e-01  7.521e-01  -0.385 0.700011
## ColorORANGE       2.922e-01  4.655e-01   0.628 0.530251
## ColorOTHER       -1.168e+00  6.442e-01  -1.812 0.069933 .
## ColorPURPLE       1.899e-01  4.250e-01   0.447 0.655029
## ColorRED         3.374e-01  2.177e-01   1.550 0.121257
```

```
## ColorSILVER      2.850e-01  2.057e-01   1.386 0.165860
## ColorWHITE       3.409e-01  2.083e-01   1.636 0.101745
## ColorYELLOW      -2.904e-01  4.947e-01  -0.587 0.557141
## WheelTypeCovers  -1.082e-01  6.698e-02  -1.615 0.106304
## WheelTypeNULL     3.489e+00  1.727e-01  20.202 < 2e-16 ***
## WheelTypeSpecial -5.363e-02  2.663e-01  -0.201 0.840390
## Odo               1.484e-05  2.184e-06   6.796 1.08e-11 ***
## SizeCROSSOVER    -9.331e-01  2.220e-01  -4.203 2.63e-05 ***
## SizeLARGE        -7.613e-01  1.319e-01  -5.770 7.91e-09 ***
## SizeLARGESUV     -7.972e-01  2.454e-01  -3.249 0.001157 **
## SizeLARGETRUCK   -1.013e+00  1.827e-01  -5.547 2.90e-08 ***
## SizeMEDIUM       -5.260e-01  1.015e-01  -5.181 2.21e-07 ***
## SizeMEDIUMSUV    -5.453e-01  1.425e-01  -3.826 0.000130 ***
## SizeSMALLSUV     -6.989e-01  2.079e-01  -3.361 0.000776 ***
## SizeSMALLTRUCK   -7.329e-01  2.520e-01  -2.908 0.003632 **
## SizeSPECIALTY    -4.271e-01  2.274e-01  -1.878 0.060352 .
## SizeSPORTS       -5.701e-01  2.545e-01  -2.240 0.025066 *
## SizeVAN          -5.982e-01  1.362e-01  -4.394 1.11e-05 ***
## MMRAuction        2.895e-05  3.634e-05   0.797 0.425670
## MMRAretail       -8.784e-06  2.241e-05  -0.392 0.695044
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 9066.3  on 6539  degrees of freedom
## Residual deviance: 7528.1  on 6480  degrees of freedom
## AIC: 7648.1
##
## Number of Fisher Scoring iterations: 5
```

4.d.

```
fitemonglm <-
  glm(formula = BadBuy ~ . ,family= 'binomial', data = dfcTrain1) %>%
  predict(dfcTest1, type = 'response') %>%
  bind_cols(dfcTest1, predictedProb=.) %>%
  mutate(predictedclass = as.factor(ifelse(predictedProb>0.5,1,0)))

fitemonglm %>%
  conf_mat(truth = BadBuy ,estimate = predictedclass) %>%
  autoplot(type = 'heatmap')
```



4.e.

```
result1 = data.frame(Auction="ADESA", Age=1, Make="HONDA", Color="SILVER",
WheelType="Covers", Odo=10000, Size="LARGE", MMRAauction=8000,
MMRAretail=10000)
```

```
predict(fitlemon, result1, type="response")
```

```
##          1
## 0.04152115
```

5.a.

```
set.seed(123)
```

```
ctrl <- trainControl(method = "repeatedcv", repeats = 10)
```

```
caretLDAresults <-
```

```
  train(BadBuy~., family = 'binomial', method = 'lda', data = dfcTrain1,
trControl = ctrl) %>%
  predict(dfcTest1, type = 'raw') %>%
  bind_cols(dfcTest1, predictedProb=.)
```

```
caretLDAresults
```

```
## # A tibble: 3,521 x 11
##   Auction Age Make Color WheelType Odo Size MMRAauction MMRAretail
BadBuy
##   <chr>    <dbl> <chr> <chr> <chr>    <dbl> <chr>    <dbl>    <dbl>
<fct>
## 1 MANHEIM      6 SATU~ WHITE Covers    81116 MEDI~    2667    3380
0
## 2 OTHER        5 CHEV~ RED Alloy    54718 MEDI~    6921    7975
1
## 3 OTHER        5 CHEV~ GOLD Covers    89365 VAN    6131    9793
1
## 4 ADESA        3 CHEV~ WHITE Covers    71794 VAN    6394    7406
0
## 5 OTHER        3 CHEV~ WHITE NULL    67229 COMP~    5785    9834
1
## 6 MANHEIM      3 DODGE GOLD Covers    71079 MEDI~    4297    5141
1
## 7 MANHEIM      6 OLDS~ SILV~ Alloy    71235 MEDI~    3325    4091
1
## 8 MANHEIM      8 PONT~ SILV~ Alloy    90325 MEDI~    2150    4937
1
## 9 MANHEIM      6 PONT~ GREEN Alloy    96893 MEDI~    4059    4884
1
## 10 OTHER       2 DODGE BLUE Covers    45151 MEDI~    7982    9121
1
## # ... with 3,511 more rows, and 1 more variable: predictedProb <fct>

caretLDAResults %>%
  xtabs(~predictedProb+BadBuy,.) %>%
  confusionMatrix(positive='1')

## Confusion Matrix and Statistics
##
##               BadBuy
## predictedProb    0    1
##               0 1377  749
##               1  405  990
##
##               Accuracy : 0.6723
##               95% CI : (0.6565, 0.6878)
##               No Information Rate : 0.5061
##               P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.3428
##
##   Mcnemar's Test P-Value : < 2.2e-16
##
##               Sensitivity : 0.5693
##               Specificity : 0.7727
##               Pos Pred Value : 0.7097
```



```
##          Neg Pred Value : 0.6477
##          Prevalence : 0.4939
##          Detection Rate : 0.2812
## Detection Prevalence : 0.3962
##          Balanced Accuracy : 0.6710
##
##          'Positive' Class : 1
##
```

```
summary(caretLDResults)
```

```
##      Auction           Age           Make           Color
## Length:3521      Min.   :1.000 Length:3521 Length:3521
## Class :character 1st Qu.:3.000 Class :character Class :character
## Mode  :character Median :4.000 Mode  :character Mode  :character
##                  Mean   :4.542
##                  3rd Qu.:6.000
##                  Max.   :9.000
##      WheelType      Odo           Size      MMRAuction
## Length:3521      Min.   : 9446 Length:3521      Min.   : 0
## Class :character 1st Qu.: 63971 Class :character 1st Qu.: 3883
## Mode  :character Median : 75523 Mode  :character Median : 5626
##                  Mean   : 73158
##                  3rd Qu.: 84057
##                  Max.   :109348
##      MMRAretail      BadBuy      predictedProb
## Min.   : 0      0:1782      0:2126
## 1st Qu.: 5916      1:1739      1:1395
## Median : 8072
## Mean   : 8219
## 3rd Qu.:10382
## Max.   :35330
```

5.b.

```
set.seed(123)
```

```
ctrl <- trainControl(method = "cv", number=10)
caretKnnresults <- train(BadBuy~., data = dfcTrain1, method = "knn",
trControl=ctrl, preProcess = c("center", "scale"), tuneLength = 10)
```


```
caretKnnresults
```

```
## k-Nearest Neighbors
##
## 6540 samples
## 9 predictor
## 2 classes: '0', '1'
##
```

```
## Pre-processing: centered (59), scaled (59)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 5885, 5886, 5887, 5887, 5885, 5886, ...
## Resampling results across tuning parameters:
##
##   k  Accuracy  Kappa
##   5  0.6079622  0.2157892
##   7  0.6098006  0.2194493
##   9  0.6096444  0.2191307
##  11  0.6131531  0.2261576
##  13  0.6100987  0.2200507
##  15  0.6140707  0.2279765
##  17  0.6197345  0.2392840
##  19  0.6211079  0.2420234
##  21  0.6189621  0.2377276
##  23  0.6174393  0.2346678
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 19.
```

5.b.i&ii

```
caret.knn.results %>%
  ggplot(aes(x=k, y=Accuracy))+ggtitle("k vs cross validation accuracy")
```



5.b.iii

```

knnresults<-caret::knnresults%>%
  predict(dfcTest1, type = 'raw') %>%
  bind_cols(dfcTest1, predictedProb=.)

knnresults %>%
  xtabs(~predictedProb+BadBuy,.) %>%
  confusionMatrix(positive='1')

## Confusion Matrix and Statistics
##
##               BadBuy
## predictedProb    0    1
##               0 1249  774
##               1  533  965
##
##               Accuracy : 0.6288
##               95% CI : (0.6126, 0.6448)
##               No Information Rate : 0.5061
##               P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.2562
##
##   Mcnemar's Test P-Value : 3.168e-11
##
##               Sensitivity : 0.5549
##               Specificity : 0.7009
##               Pos Pred Value : 0.6442
##               Neg Pred Value : 0.6174
##               Prevalence : 0.4939
##               Detection Rate : 0.2741
##               Detection Prevalence : 0.4254
##               Balanced Accuracy : 0.6279
##
##               'Positive' Class : 1
##

```

5.c.

```

library("glmnet")

## Warning: package 'glmnet' was built under R version 3.6.3

## Loading required package: Matrix

##
## Attaching package: 'Matrix'

## The following objects are masked from 'package:tidyr':
##
##   expand, pack, unpack

```

```
## Loaded glmnet 3.0-2

lambdaValues <- 10^seq(-5, 2, length = 100)

set.seed(123)


fitlemonlasso <- train(BadBuy ~ ., family='binomial', data=dfcTrain1,
method='glmnet', trControl=trainControl(method='cv', number=10), tuneGrid =
expand.grid(alpha=1, lambda=lambdaValues))

varImp(fitlemonlasso)$importance %>%      # Add scale=FALSE inside VarImp if
you don't want to scale
  rownames_to_column(var = "Variable") %>%
  mutate(Importance = scales::percent(Overall/100)) %>%
  arrange(desc(Overall)) %>%
  as_tibble()

## # A tibble: 59 x 3
##   Variable      Overall Importance
##   <chr>         <dbl> <chr>
## 1 WheelTypeNULL    100  100%
## 2 ColorOTHER       36.2  36%
## 3 SizeLARGETRUCK   26.5  26%
## 4 SizeCROSSOVER    24.2  24%
## 5 SizeLARGE        20.1  20%
## 6 MakeLINCOLN      19.7  20%
## 7 SizeLARGESUV     19.6  20%
## 8 MakeSUZUKI       19.4  19%
## 9 SizeSMALLTRUCK   18.8  19%
## 10 MakeMITSUBISHI  18.1  18%
## # ... with 49 more rows
```

5.c.ii

```
plot(varImp(fitlemonlasso), top = 25)
```



5.c.iii

```
fitlemonlasso$bestTune$lambda
```

```
## [1] 0.0003053856
```

5.c.iv.

```
Lassoresults <-
  fitlemonlasso %>%
  predict(dfctest1, type='raw') %>%
  bind_cols(dfctest1, predictedClass=.)

Lassoresults %>%
  xtabs(~predictedClass+BadBuy, .) %>%
  confusionMatrix(positive = '1')

## Confusion Matrix and Statistics
##
##               BadBuy
## predictedClass  0    1
## 0      1339    721
## 1       443   1018
##
##               Accuracy : 0.6694
##               95% CI   : (0.6536, 0.6849)
##      No Information Rate : 0.5061
##      P-Value [Acc > NIR] : < 2e-16
```

```
##
##           Kappa : 0.3374
##
## McNemar's Test P-Value : 4.7e-16
##
##           Sensitivity : 0.5854
##           Specificity : 0.7514
##           Pos Pred Value : 0.6968
##           Neg Pred Value : 0.6500
##           Prevalence : 0.4939
##           Detection Rate : 0.2891
##           Detection Prevalence : 0.4149
##           Balanced Accuracy : 0.6684
##
##           'Positive' Class : 1
##
```

5.d.i

```
set.seed(123)

fitlemonridge <- train(BadBuy ~ ., family='binomial', data=dfcTrain1,
method='glmnet', trControl=trainControl(method='cv', number=10), tuneGrid =
expand.grid(alpha=0, lambda=lambdaValues))

ridgeresults <-
  fitlemonridge %>%
  predict(dfcTest1, type='raw') %>%
  bind_cols(dfcTest1, predictedClass=.)

ridgeresults %>%
  xtabs(~predictedClass+BadBuy, .) %>%
  confusionMatrix(positive = '1')

## Confusion Matrix and Statistics
##
##           BadBuy
## predictedClass  0    1
##           0 1323  699
##           1  459 1040
##
##           Accuracy : 0.6711
##           95% CI : (0.6553, 0.6866)
##           No Information Rate : 0.5061
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.341
##
## McNemar's Test P-Value : 2.166e-12
##
##           Sensitivity : 0.5980
```

```
##           Specificity : 0.7424
##           Pos Pred Value : 0.6938
##           Neg Pred Value : 0.6543
##           Prevalence : 0.4939
##           Detection Rate : 0.2954
##           Detection Prevalence : 0.4257
##           Balanced Accuracy : 0.6702
##
##           'Positive' Class : 1
##
```

```
fitLemonridge$bestTune$lambda
```

```
## [1] 0.0559081
```

5.d.ii

```
#elastic net
set.seed(123)
fitLemonElastic <- train(BadBuy ~ ., family='binomial', data=dfcTrain1,
method='glmnet', trControl=trainControl(method='cv', number=10),
tuneLength=10)
```

```
elasticResults <-
  fitLemonElastic %>%
  predict(dfcTest1, type='raw') %>%
  bind_cols(dfcTest1, predictedClass=.)
```

```
elasticResults %>%
  xtabs(~predictedClass+BadBuy, .) %>%
  confusionMatrix(positive = '1')
```

```
## Confusion Matrix and Statistics
```

```
##
##           BadBuy
## predictedClass  0    1
##           0 1338  721
##           1  444 1018
##
##           Accuracy : 0.6691
##           95% CI : (0.6533, 0.6847)
##           No Information Rate : 0.5061
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.3369
##
##           Mcnemar's Test P-Value : 6.154e-16
##
##           Sensitivity : 0.5854
##           Specificity : 0.7508
##           Pos Pred Value : 0.6963
```

```
##           Neg Pred Value : 0.6498
##           Prevalence : 0.4939
##           Detection Rate : 0.2891
## Detection Prevalence : 0.4152
##           Balanced Accuracy : 0.6681
##
##           'Positive' Class : 1
##
```

5.e

```
set.seed(123)

fitlemonqda <- train(BadBuy ~ ., family= "binomial", data=dfcTrain1,
method="qda", trControl = trainControl(method='cv', number=10))

## Warning: model fit failed for Fold03: parameter=none Error in
qda.default(x, grouping, ...) : rank deficiency in group 0

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info =
trainInfo, :
## There were missing values in resampled performance measures.

qdaresults <- fitlemonqda %>%
  predict(dfcTest1, type = 'raw') %>%
  bind_cols(dfcTest1, predictedProb=.)

qdaresults

## # A tibble: 3,521 x 11
##   Auction Age Make Color WheelType Odo Size MMRAauction MMRAretail
BadBuy
##   <chr> <dbl> <chr> <chr> <chr> <dbl> <chr> <dbl> <dbl>
<fct>
## 1 MANHEIM 6 SATU~ WHITE Covers 81116 MEDI~ 2667 3380
0
## 2 OTHER 5 CHEV~ RED Alloy 54718 MEDI~ 6921 7975
1
## 3 OTHER 5 CHEV~ GOLD Covers 89365 VAN 6131 9793
1
## 4 ADESA 3 CHEV~ WHITE Covers 71794 VAN 6394 7406
0
## 5 OTHER 3 CHEV~ WHITE NULL 67229 COMP~ 5785 9834
1
## 6 MANHEIM 3 DODGE GOLD Covers 71079 MEDI~ 4297 5141
1
## 7 MANHEIM 6 OLDS~ SILV~ Alloy 71235 MEDI~ 3325 4091
1
## 8 MANHEIM 8 PONT~ SILV~ Alloy 90325 MEDI~ 2150 4937
1
## 9 MANHEIM 6 PONT~ GREEN Alloy 96893 MEDI~ 4059 4884
```



```

1
## 10 OTHER          2 DODGE BLUE  Covers    45151 MEDI~          7982          9121
1
## # ... with 3,511 more rows, and 1 more variable: predictedProb <fct>

qdaresults %>%
  xtabs(~predictedProb+BadBuy, .) %>%
  confusionMatrix(positive='1')

## Confusion Matrix and Statistics
##
##              BadBuy
## predictedProb    0    1
##      0 1483   973
##      1   299   766
##
##              Accuracy : 0.6387
##              95% CI   : (0.6226, 0.6546)
##      No Information Rate : 0.5061
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa   : 0.274
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.4405
##              Specificity : 0.8322
##              Pos Pred Value : 0.7192
##              Neg Pred Value : 0.6038
##              Prevalence : 0.4939
##              Detection Rate : 0.2176
##      Detection Prevalence : 0.3025
##              Balanced Accuracy : 0.6363
##
##              'Positive' Class : 1
##

```